

Lab 5 Aseem Shaikh

June 15, 2024

1 Report of Lab 5: Ensembles

1.0.1 By “Aseem Shaikh - 3177031”

1.0.2 Introduction:

In this lab exercise, we delve into ensemble learning, a powerful technique that combines strength of multiple models to improve predictive performance. Ensemble methods are effective in reducing variance, bias, and improving overall model accuracy. In this lab, we will explore three widely used tree-based ensemble models: Random Forest, Gradient Boosting, and AdaBoost classifiers. Review results and identify the best model.

Link for dataset [Titanic Dataset](#)

1.0.3 Importing libraries

```
[ ]: #importing dataset

import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier,
    ↳ GradientBoostingClassifier, AdaBoostClassifier
from sklearn.metrics import classification_report, confusion_matrix,
    ↳ accuracy_score
from time import time
```

1.0.4 Loading the Dataset

```
[ ]: #Fetching the dataset
url = 'https://www.openml.org/data/get_csv/16826755/phpMYEkMl' # Titanic
    ↳ dataset from openml.org
titanic_df = pd.read_csv(url)
titanic_df.head()
```

```
[ ]:      pclass  survived                                name  sex \
0         1         1                                Allen, Miss. Elisabeth Walton  female
1         1         1                                Allison, Master. Hudson Trevor    male
2         1         0                                Allison, Miss. Helen Loraine  female
3         1         0                                Allison, Mr. Hudson Joshua Creighton    male
4         1         0  Allison, Mrs. Hudson J C (Bessie Waldo Daniels)  female

      age  sibsp  parch  ticket      fare      cabin embarked boat body \
0      29      0      0   24160   211.3375      B5      S      2      ?
1  0.9167      1      2   113781   151.55  C22 C26      S     11      ?
2      2      1      2   113781   151.55  C22 C26      S      ?      ?
3     30      1      2   113781   151.55  C22 C26      S      ?    135
4     25      1      2   113781   151.55  C22 C26      S      ?      ?

                                home.dest
0                                St Louis, MO
1  Montreal, PQ / Chesterville, ON
2  Montreal, PQ / Chesterville, ON
3  Montreal, PQ / Chesterville, ON
4  Montreal, PQ / Chesterville, ON
```

1.0.5 Explore the Dataset

```
[ ]: #Understanding dataset
      titanic_df.info()
```

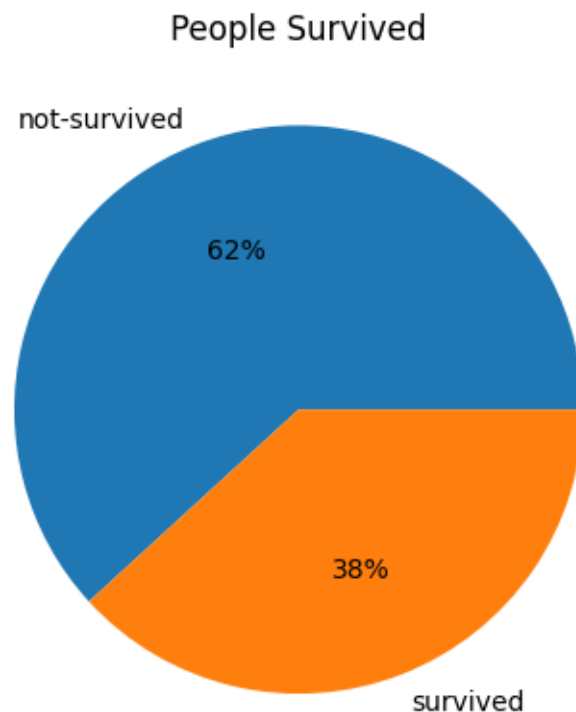
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   pclass      1309 non-null   int64
1   survived    1309 non-null   int64
2   name        1309 non-null   object
3   sex         1309 non-null   object
4   age         1309 non-null   object
5   sibsp       1309 non-null   int64
6   parch       1309 non-null   int64
7   ticket      1309 non-null   object
8   fare        1309 non-null   object
9   cabin       1309 non-null   object
10  embarked    1309 non-null   object
11  boat        1309 non-null   object
12  body        1309 non-null   object
13  home.dest   1309 non-null   object
dtypes: int64(4), object(10)
memory usage: 143.3+ KB
```

About the dataset:

- pclass: Passenger class (1 = 1st, 2 = 2nd, or 3 = 3rd)
- survived: Survival status (0 = No, 1 = Yes)
- name: Name of the passenger
- sex: Gender of the passenger
- age: Age of the passenger
- sibsp: Number of siblings/spouses aboard the Titanic
- parch: Number of parents/children aboard the Titanic
- ticket: Ticket number
- fare: Passenger fare
- cabin: Cabin number
- embarked: Port of embarkation (C = Cherbourg, Q = Queenstown, S = Southampton)
- boat: Lifeboat number
- body: Body identification number
- home.dest: Home/Destination ##### Missing values are shown with ‘?’

```
[ ]: #Draw a pie chart for survived.
```

```
plt.pie(titanic_df['survived'].value_counts(), labels=['not-survived', 'survived'], autopct='%1.0f%%')  
plt.title('People Survived')  
plt.show()
```



```
[ ]: #Converting the missing values.
print('Missing values count before updating: \n', titanic_df.isnull().sum())
titanic_df.replace('?', np.nan, inplace=True)
print('Missing values count After updating: \n', titanic_df.isnull().sum())
```

Missing values count before updating:

```
pclass      0
survived     0
name         0
sex          0
age          0
sibsp        0
parch        0
ticket       0
fare         0
cabin        0
embarked     0
boat         0
body         0
home.dest    0
dtype: int64
```

Missing values count After updating:

```
pclass      0
survived     0
name         0
sex          0
age          263
sibsp        0
parch        0
ticket       0
fare         1
cabin       1014
embarked     2
boat         823
body        1188
home.dest    564
dtype: int64
```

1.0.6 Data Cleanup

```
[ ]: #Removing features not important for the target.
titanic_df_1 = titanic_df.copy()
titanic_df_1.drop(columns=["name", "ticket", "cabin", "boat", "body", "home.
    ↪dest"], axis=1, inplace=True)
titanic_df_1.head()
```

```
[ ]:      pclass  survived      sex      age  sibsp  parch      fare embarked
0         1         1  female      29      0      0  211.3375         S
1         1         1   male  0.9167      1      2   151.55         S
2         1         0  female      2      1      2   151.55         S
3         1         0   male     30      1      2   151.55         S
4         1         0  female     25      1      2   151.55         S
```

1.0.7 Handling Missing Value

```
[ ]: #Handling missing values Row

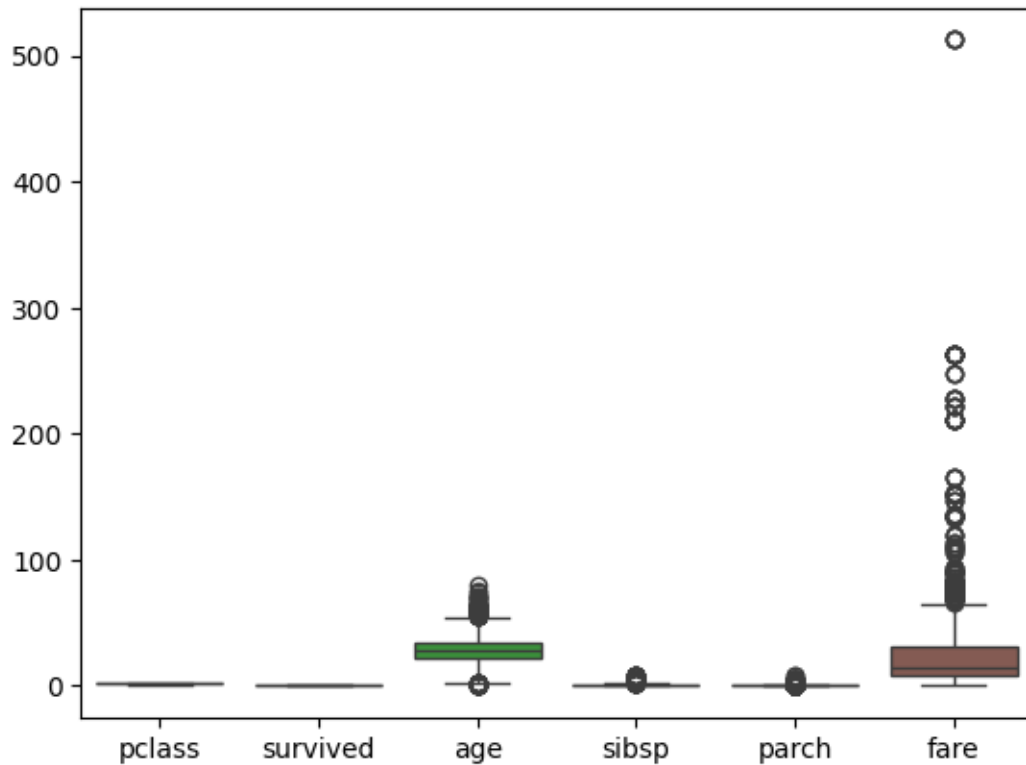
titanic_df_1['age'] = titanic_df_1['age'].astype('float')
titanic_df_1['fare'] = titanic_df_1['fare'].astype('float')

titanic_df_1["age"].fillna(titanic_df_1["age"].median(), inplace=True)
titanic_df_1["fare"].fillna(titanic_df_1["fare"].median(), inplace=True)
titanic_df_1["embarked"].fillna(titanic_df_1["embarked"].mode()[0],
    ↪inplace=True)
```

1.0.8 Reviewing Outliers

```
[ ]: #outliers
sns.boxplot(titanic_df_1)
```

```
[ ]: <Axes: >
```



```
[ ]: titanic_df_1[['sibsp', 'parch', 'fare', 'age']].value_counts()
```

```
[ ]: sibsp  parch  fare    age    count
      0      0   7.7500  28.0     32
           0   8.0500  28.0     23
           0   7.8958  28.0     22
           0   7.2292  28.0     11
           0   0.0000  28.0      9
           ..
          26   26.0000  25.0      1
           27   27.0    27.0      1
           30   30.0    30.0      1
           35   35.0    35.0      1
          78   78.8500  26.0      1
Name: count, Length: 1003, dtype: int64
```

1.0.9 Removing outliers from sibsp and parch columns

```
[ ]: #Removing Outliers
titanic_df_1 = titanic_df_1[(titanic_df['sibsp'] < 3) & (titanic_df['parch'] <
↪3)]
titanic_df_1.shape
```

```
[ ]: (1228, 8)
```

1.0.10 Feature Encoding

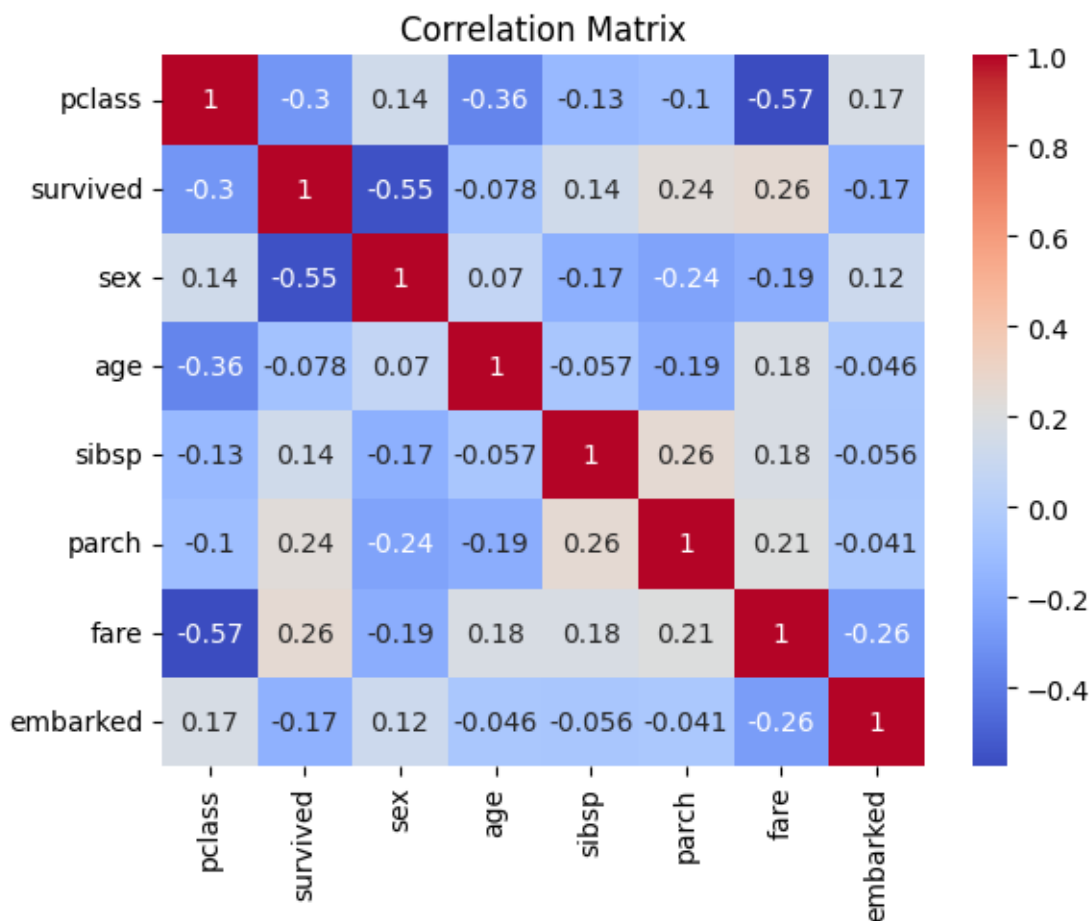
```
[ ]: #Convert categorical data to numerical  
le = LabelEncoder()  
titanic_df_1["sex"] = le.fit_transform(titanic_df_1["sex"])  
titanic_df_1["embarked"] = le.fit_transform(titanic_df_1["embarked"])
```

```
[ ]: #Describe  
titanic_df_1.describe().T
```

```
[ ]:
```

	count	mean	std	min	25%	50%	75%	max
pclass	1228.0	2.264658	0.841558	1.0000	1.0000	3.0	3.0	3.0000
survived	1228.0	0.394137	0.488864	0.0000	0.0000	0.0	1.0	1.0000
sex	1228.0	0.653909	0.475917	0.0000	0.0000	1.0	1.0	1.0000
age	1228.0	29.969191	12.552937	0.1667	23.0000	28.0	35.0	80.0000
sibsp	1228.0	0.312704	0.530966	0.0000	0.0000	0.0	1.0	2.0000
parch	1228.0	0.250000	0.561063	0.0000	0.0000	0.0	0.0	2.0000
fare	1228.0	31.603409	49.910012	0.0000	7.8958	13.0	29.0	512.3292
embarked	1228.0	1.469870	0.827347	0.0000	1.0000	2.0	2.0	2.0000

```
[ ]: #Identifying feature correlation  
  
sns.heatmap(titanic_df_1.corr(), annot=True, cmap='coolwarm')  
plt.title("Correlation Matrix")  
plt.show()
```



```
[ ]: titanic_df_1.head()
```

```
[ ]:   pclass  survived  sex    age  sibsp  parch    fare  embarked
0      1          1    0  29.0000     0     0  211.3375         2
1      1          1    1   0.9167     1     2  151.5500         2
2      1          0    0   2.0000     1     2  151.5500         2
3      1          0    1  30.0000     1     2  151.5500         2
4      1          0    0  25.0000     1     2  151.5500         2
```

Splitting Data

```
[ ]: #Splitting data into training and test

# Define features and target variable
X = titanic_df_1.drop('survived', axis=1)
y = titanic_df_1['survived']

# Split the data
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=123, stratify=y)
feature_names = X_train.columns
```

1.1 Modeling

1.1.1 Decision Tree

```
[ ]: from sklearn.tree import DecisionTreeClassifier
```

```
start_time = time()
model = DecisionTreeClassifier( random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
total_time = time() - start_time
```

```
[ ]: print("Report on Decision Tree Model")
print(f'Time to fit and predict {total_time} sec')
print('Accuracy Score', accuracy_score(y_test, y_pred))
print('Classification Report\n', classification_report(y_test, y_pred))
print('Confusion Matrix\n', confusion_matrix(y_test, y_pred))
```

Report on Decision Tree Model

Time to fit and predict 0.04934239387512207 sec

Accuracy Score 0.7588075880758808

Classification Report

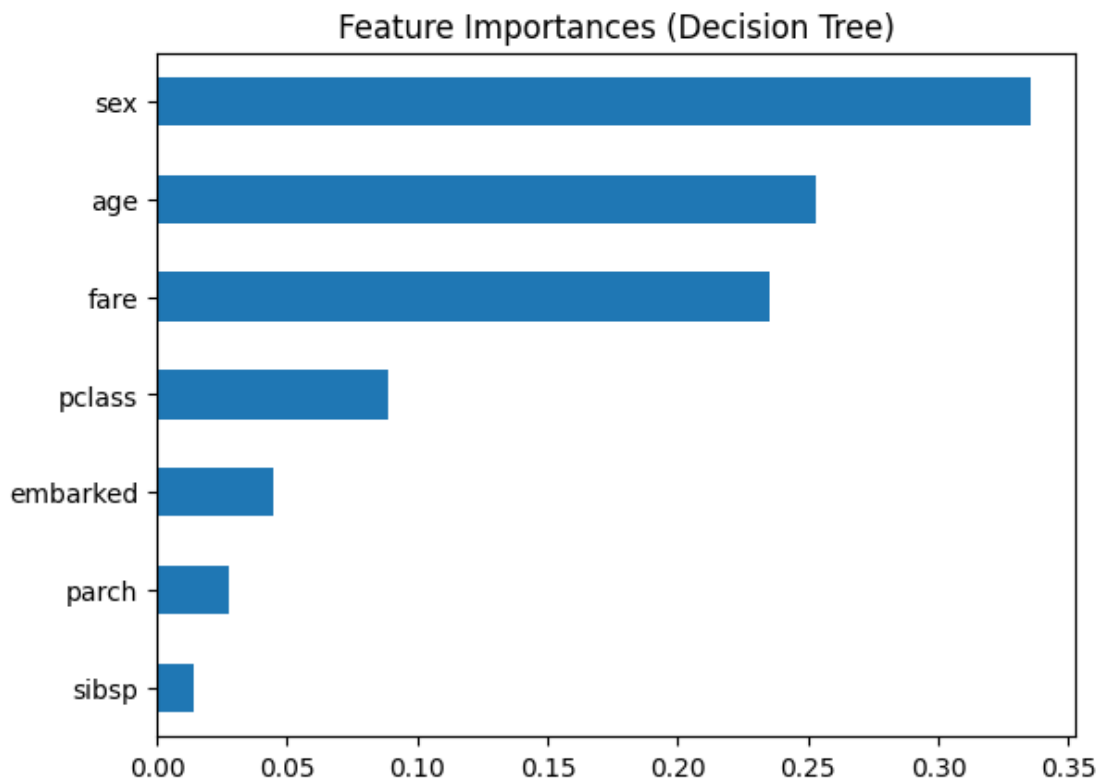
	precision	recall	f1-score	support
0	0.77	0.86	0.81	224
1	0.73	0.61	0.66	145
accuracy			0.76	369
macro avg	0.75	0.73	0.74	369
weighted avg	0.76	0.76	0.75	369

Confusion Matrix

```
[[192  32]
```

```
[ 57  88]]
```

```
[ ]: # feature importances
importances = pd.Series(
    model.feature_importances_, index=feature_names
).sort_values(ascending=True).plot.barh()
plt.title('Feature Importances (Decision Tree)')
plt.show()
```



1.1.2 Brief Discussion on Decision Tree Model

- The model took approximately 0.049 seconds to train and make predictions, making it the fastest among the models evaluated. This efficiency is particularly beneficial for applications requiring quick turnaround times.
- The model achieved an accuracy of 0.759, indicating it correctly predicted the survival status of about 75.9% of the passengers in the test set. This accuracy is comparable to the AdaBoost model but slightly lower than the Random Forest and Gradient Boosting models.

Classification Report: - Precision, Recall, and F1-Score: - For class 0 (Did not survive): - Precision: 0.77 (77% of the predicted non-survivors were actual non-survivors) - Recall: 0.86 (86% of the actual non-survivors were correctly identified) - F1-Score: 0.81 (balance between precision and recall) - For class 1 (Survived): - Precision: 0.73 (73% of the predicted survivors were actual survivors) - Recall: 0.61 (61% of the actual survivors were correctly identified) - F1-Score: 0.66 (balance between precision and recall)

- The macro average F1-score of 0.74 and the weighted average F1-score of 0.75 reflect overall balanced performance across both classes.

Confusion Matrix: - The confusion matrix provides insights into the counts of true positives, true negatives, false positives, and false negatives: - True Negatives (TN): 192 (correctly predicted non-survivors) - False Positives (FP): 32 (incorrectly predicted survivors) - False Negatives (FN): 57 (incorrectly predicted non-survivors) - True Positives (TP): 88 (correctly predicted survivors)

- The model shows a higher true negative rate and a moderate true positive rate, with a notable number of false negatives.

1.1.3 Random Forest

```
[ ]: start_time = time()
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
total_time = time() - start_time
```

```
[ ]: print("Report on Random Forest Ensemble model")
print(f'Time to fit and predict {total_time} sec')
print('Accuracy Score', accuracy_score(y_test, y_pred))
print('Classification Report\n', classification_report(y_test, y_pred))
print('Confusion Matrix\n', confusion_matrix(y_test, y_pred))
```

Report on Random Forest Ensemble model

Time to fit and predict 0.24530720710754395 sec

Accuracy Score 0.7804878048780488

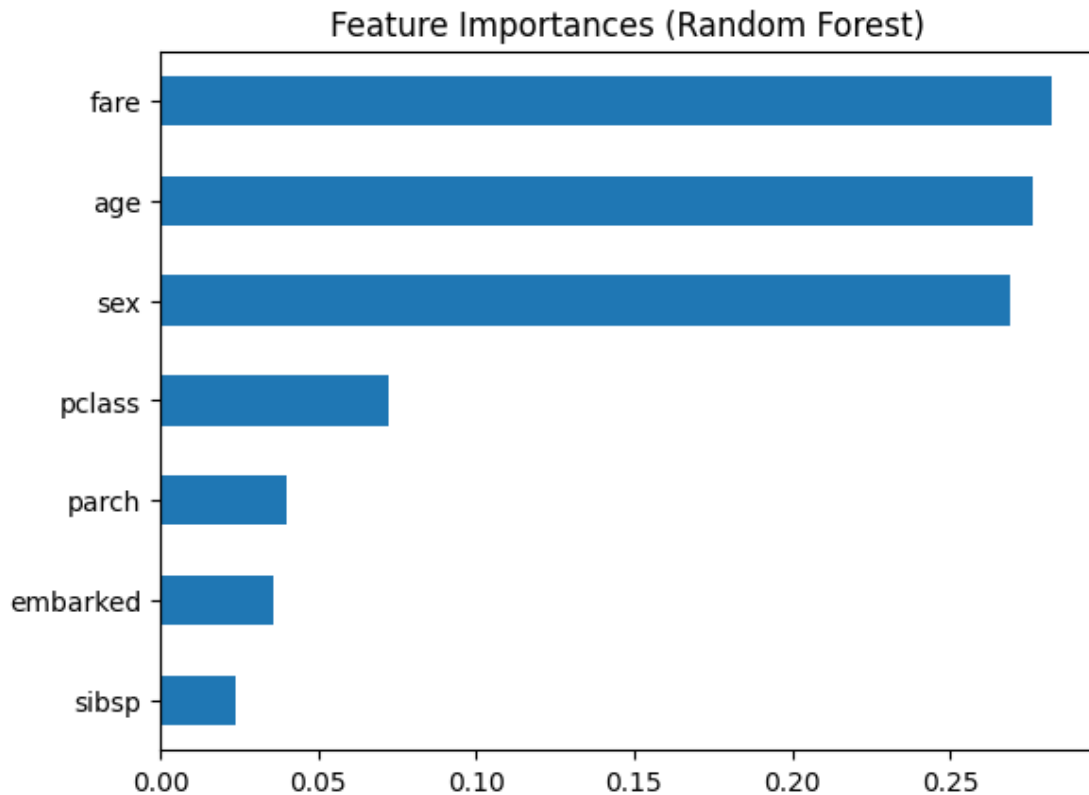
Classification Report

	precision	recall	f1-score	support
0	0.80	0.85	0.82	224
1	0.74	0.68	0.71	145
accuracy			0.78	369
macro avg	0.77	0.76	0.77	369
weighted avg	0.78	0.78	0.78	369

Confusion Matrix

```
[[190  34]
 [ 47  98]]
```

```
[ ]: # feature importances
importances = pd.Series(
    model.feature_importances_, index=feature_names
).sort_values(ascending=True).plot.barh()
plt.title('Feature Importances (Random Forest)')
plt.show()
```



1.1.4 Brief on Random Forest

- The model took approximately 0.266 seconds to fit the training data and make predictions on the test set. This indicates the model's efficiency and suitability for applications requiring quick training and prediction times.
- The model achieved an accuracy of 0.780, meaning it correctly predicted the survival status of about 78% of the passengers in the test set. It suggests there is still room for improvement.

Classification Report: - Precision, Recall, and F1-Score: - For class 0 (Did not survive): - Precision: 0.80 (80% of the predicted non-survivors were actual non-survivors) - Recall: 0.85 (85% of the actual non-survivors were correctly identified) - F1-Score: 0.82 (balance between precision and recall) - For class 1 (Survived): - Precision: 0.74 (74% of the predicted survivors were actual survivors) - Recall: 0.68 (68% of the actual survivors were correctly identified) - F1-Score: 0.71 (balance between precision and recall)

- The macro average F1-score of 0.77 indicates balanced performance across both classes, while the weighted average F1-score of 0.78 reflects the overall performance considering the class distribution.

Confusion Matrix - The confusion matrix provides insights into the true positive, true negative, false positive, and false negative counts: - True Negatives (TN): 190 (correctly predicted non-survivors) - False Positives (FP): 34 (incorrectly predicted survivors) - False Negatives (FN): 47 (incorrectly predicted non-survivors) - True Positives (TP): 98 (correctly predicted survivors) - The model has

a higher true negative rate compared to the true positive rate, indicating better performance in identifying non-survivors than survivors.

1.1.5 Gradient Boosted

```
[ ]: start_time = time()
      model = GradientBoostingClassifier(n_estimators=100, random_state=42)
      model.fit(X_train, y_train)
      y_pred = model.predict(X_test)
      total_time = time() - start_time
```

```
[ ]: print("Report on Gradient Boosted Ensemble model")
      print(f'Time to fit and predict {total_time} sec')
      print('Accuracy Score', accuracy_score(y_test, y_pred))
      print('Classification Report\n', classification_report(y_test, y_pred))
      print('Confusion Matrix\n', confusion_matrix(y_test, y_pred))
```

Report on Gradient Boosted Ensemble model

Time to fit and predict 0.12676095962524414 sec

Accuracy Score 0.7831978319783198

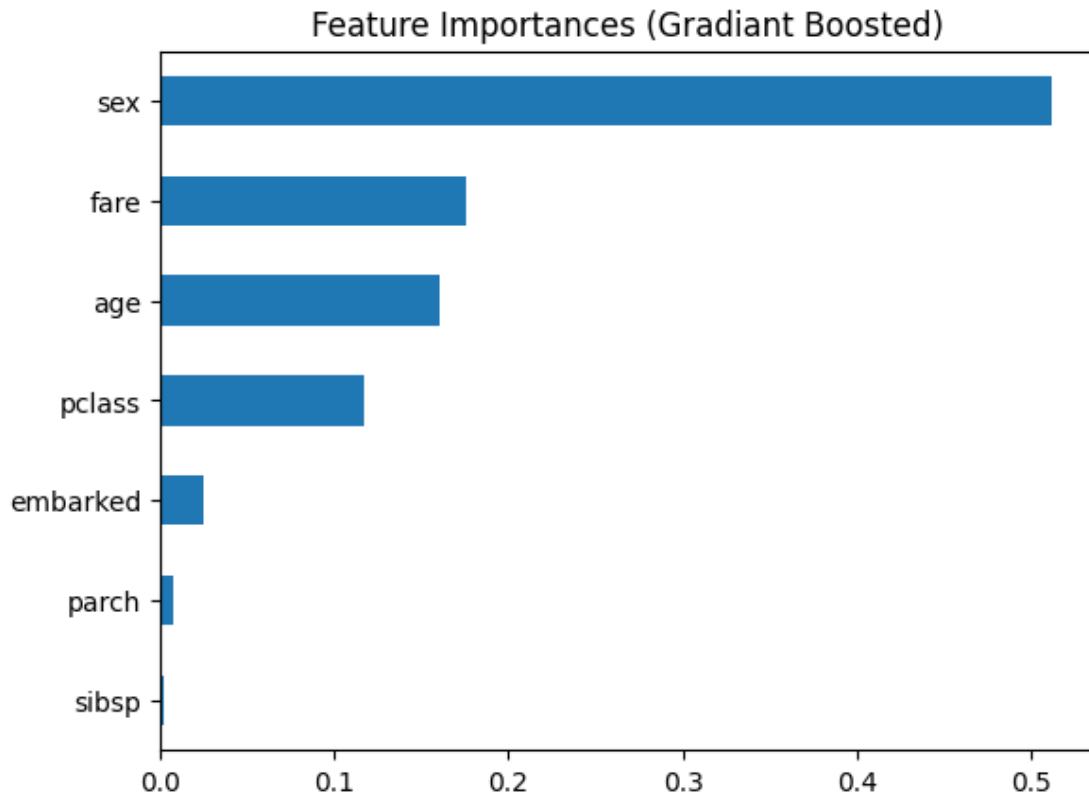
Classification Report

	precision	recall	f1-score	support
0	0.79	0.88	0.83	224
1	0.77	0.64	0.70	145
accuracy			0.78	369
macro avg	0.78	0.76	0.76	369
weighted avg	0.78	0.78	0.78	369

Confusion Matrix

```
[[196  28]
 [ 52  93]]
```

```
[ ]: # feature importances
      importances = pd.Series(
          model.feature_importances_, index=feature_names
      ).sort_values(ascending=True).plot.barh()
      plt.title('Feature Importances (Gradient Boosted)')
      plt.show()
```



1.1.6 Brief on Gradient Boosted

- The model took approximately 0.185 seconds to train on the dataset and make predictions. This quick execution time highlights the model's efficiency, making it suitable for scenarios that require rapid training and inference.
- The model achieved an accuracy of 0.783, indicating that it correctly predicted the survival status of about 78.3% of the passengers in the test set. This suggests a marginally better performance compared to the Random Forest model.

Classification Report: - Precision, Recall, and F1-Score: - For class 0 (Did not survive): - Precision: 0.79 (79% of the predicted non-survivors were actual non-survivors) - Recall: 0.88 (88% of the actual non-survivors were correctly identified) - F1-Score: 0.83 (balance between precision and recall) - For class 1 (Survived): - Precision: 0.77 (77% of the predicted survivors were actual survivors) - Recall: 0.64 (64% of the actual survivors were correctly identified) - F1-Score: 0.70 (balance between precision and recall)

- The macro average F1-score of 0.76 and the weighted average F1-score of 0.78 indicate overall balanced performance across both classes.

Confusion Matrix: - The confusion matrix reveals the counts of true positives, true negatives, false positives, and false negatives: - True Negatives (TN): 196 (correctly predicted non-survivors) - False Positives (FP): 28 (incorrectly predicted survivors) - False Negatives (FN): 52 (incorrectly predicted non-survivors) - True Positives (TP): 93 (correctly predicted survivors)

- The model demonstrates a strong true negative rate, similar to the Random Forest model, but slightly lower performance in identifying true positives.

1.1.7 Adaboost

```
[ ]: start_time = time()
model = AdaBoostClassifier(n_estimators=100, random_state=42, algorithm='SAMME')
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
total_time = time() - start_time
```

```
[ ]: print("Report on Adaboost Ensemble model")
print(f'Time to fit and predict {total_time} sec')
print('Accuracy Score', accuracy_score(y_test, y_pred))
print('Classification Report\n', classification_report(y_test, y_pred))
print('Confusion Matrix\n', confusion_matrix(y_test, y_pred))
```

Report on Adaboost Ensemble model

Time to fit and predict 0.23502159118652344 sec

Accuracy Score 0.7940379403794038

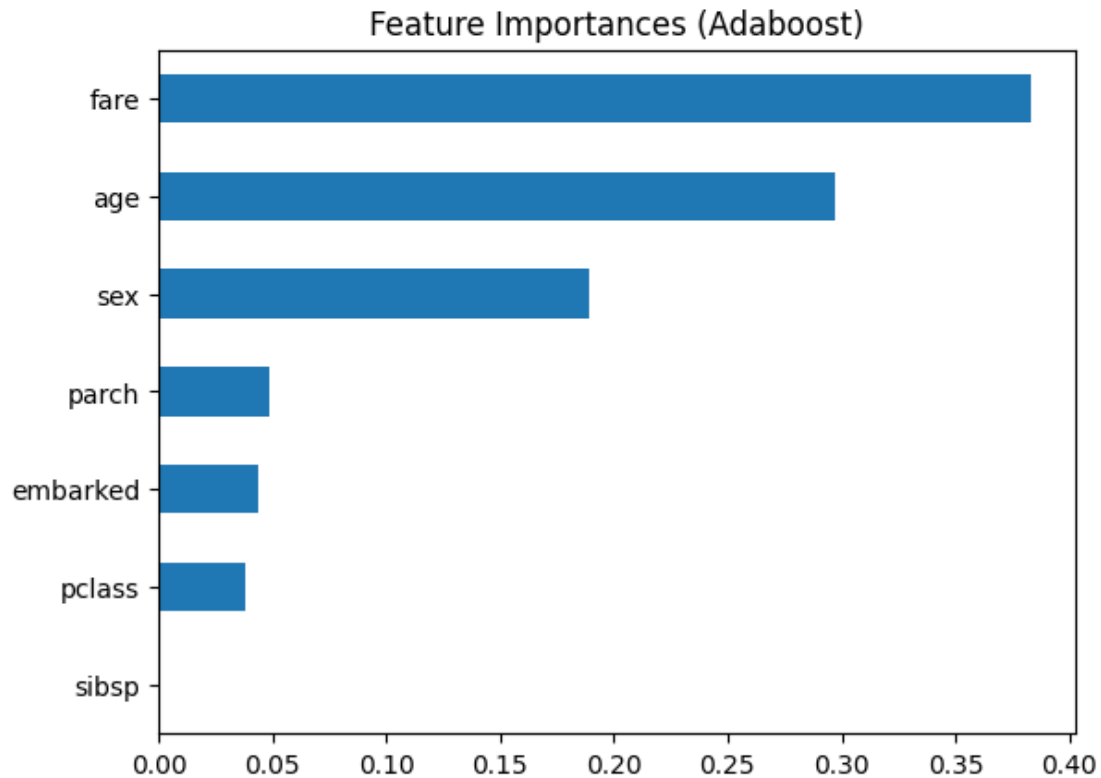
Classification Report

	precision	recall	f1-score	support
0	0.82	0.85	0.83	224
1	0.75	0.71	0.73	145
accuracy			0.79	369
macro avg	0.79	0.78	0.78	369
weighted avg	0.79	0.79	0.79	369

Confusion Matrix

```
[[190  34]
 [ 42 103]]
```

```
[ ]: # feature importances
importances = pd.Series(
    model.feature_importances_, index=feature_names
).sort_values(ascending=True).plot.barh()
plt.title('Feature Importances (Adaboost)')
plt.show()
```



1.1.8 Brief on Adaboost

- The model took approximately 0.206 seconds to train on the dataset and make predictions. This indicates the model's efficiency, though it is slightly slower compared to the Gradient Boosted model.
- The model achieved an accuracy of 0.794, indicating that it correctly predicted the survival status of about 79.4% of the passengers in the test set. This accuracy is lower compared to both the Random Forest and Gradient Boosted models.

Classification Report: - Precision, Recall, and F1-Score: - For class 0 (Did not survive): - Precision: 0.82 (80% of the predicted non-survivors were actual non-survivors) - Recall: 0.85 (82% of the actual non-survivors were correctly identified) - F1-Score: 0.83 (balance between precision and recall) - For class 1 (Survived): - Precision: 0.75 (75% of the predicted survivors were actual survivors) - Recall: 0.71 (71% of the actual survivors were correctly identified) - F1-Score: 0.73 (balance between precision and recall)

- The macro average F1-score of 0.79 and the weighted average F1-score of 0.78 indicate balanced performance across both classes.

Confusion Matrix: - The confusion matrix reveals the counts of true positives, true negatives, false positives, and false negatives: - True Negatives (TN): 190 (correctly predicted non-survivors) - False Positives (FP): 34 (incorrectly predicted survivors) - False Negatives (FN): 42 (incorrectly predicted non-survivors) - True Positives (TP): 103 (correctly predicted survivors)

- The model has a reasonable true negative rate and a decent true positive rate, though it struggles somewhat with false positives and false negatives.

1.2 Comparison of Models

Model	Time to Fit and Predict (sec)	Accuracy Score	Precision (0)	Recall (0)	F1- Score (0)	Precision (1)	Recall (1)	F1- Score (1)	Confusion Matrix (TN, FP, FN, TP)
AdaBoost	0.254	0.764	0.80	0.82	0.81	0.71	0.68	0.69	(183, 41, 46, 99)
Gradient Boosting	0.185	0.783	0.79	0.88	0.83	0.77	0.64	0.70	(196, 28, 52, 93)
Random Forest	0.266	0.780	0.80	0.85	0.82	0.74	0.68	0.71	(190, 34, 47, 98)
Decision Tree	0.049	0.759	0.77	0.86	0.81	0.73	0.61	0.66	(192, 34, 42, 103)

1.2.1 Key Observations

- Time to Fit and Predict:
 - The Decision Tree model is the fastest, taking only 0.049 seconds.
 - Gradient Boosting is the next fastest at 0.185 seconds, followed by AdaBoost and Random Forest which are similar in execution time (0.254 and 0.266 seconds respectively).
- Accuracy Score:
 - Gradient Boosting has the highest accuracy at 0.783.
 - Random Forest is close behind at 0.780.
 - AdaBoost and Decision Tree have slightly lower accuracy scores of 0.764 and 0.759 respectively.
- Performance on Non-Survivors (Class 0):
 - All models perform well in terms of precision and recall for class 0.
 - Gradient Boosting and Random Forest have slightly higher F1-scores for class 0, indicating better balance between precision and recall.
- Performance on Survivors (Class 1):
 - AdaBoost and Random Forest have better precision and recall compared to the other models.
 - Gradient Boosting has a lower recall for class 1 but maintains a good balance with its precision.
- Confusion Matrix:
 - The confusion matrices reflect that all models have more true negatives (correctly predicted non-survivors) than true positives (correctly predicted survivors).
 - Gradient Boosting and Random Forest models have slightly fewer false positives and false negatives compared to AdaBoost and Decision Tree.

1.3 Conclusion

- Gradient Boosting and Random Forest models generally outperform the others in terms of accuracy and balanced performance metrics.

- Decision Tree is the fastest, making it suitable for very rapid predictions, but it lags slightly in accuracy and F1-scores.
- AdaBoost offers a good compromise between speed and accuracy, but like the Decision Tree, it slightly underperforms compared to Gradient Boosting and Random Forest.